

<b>Introduction</b>	2
<b>1. Schema Design and Data Acquisition</b>	2
1.1 E/R diagram	2
1.2 Commands for creating tables	3
1.3 Data Acquisition	6
1.3.1 publication.csv	6
1.3.2 author.csv	6
1.3.3 pub_author.csv	7
This file has two columns: pub_id and author_name.	7
It stores a many-to-many relationship between publications and authors. For example, in Fig4, pubid 3 has 2 authors: Nathan Goodman and Oded Shmueli respectively.	7
1.3.4 Populate database tables	7
<b>2. Queries and Optimizing Queries</b>	9
2.1 SQL queries & screen capture of results	24
2.2 Figures and analysis	25
<b>3. Effect of Index</b>	27
3.1 Create index command	27
3.2 Analysis on Index	27
<b>4. Effect of Cache</b>	30
4.1 Analysis on effect of cache	30
<b>5. Appendix</b>	32
5.1 Java code SAX Parser	32

# Introduction

In this project, we have gained experience in various stages of database design: schema design, data acquisition, data transformation, querying and indexing. Moreover, we have worked with large dataset and have understood the importance of efficient SQL queries.

For this project, we have utilized MySQL, Excel, and Java as the main programming tools.

## 1. Schema Design and Data Acquisition

### 1.1 E/R diagram

The ER diagram of the schema is provided below.

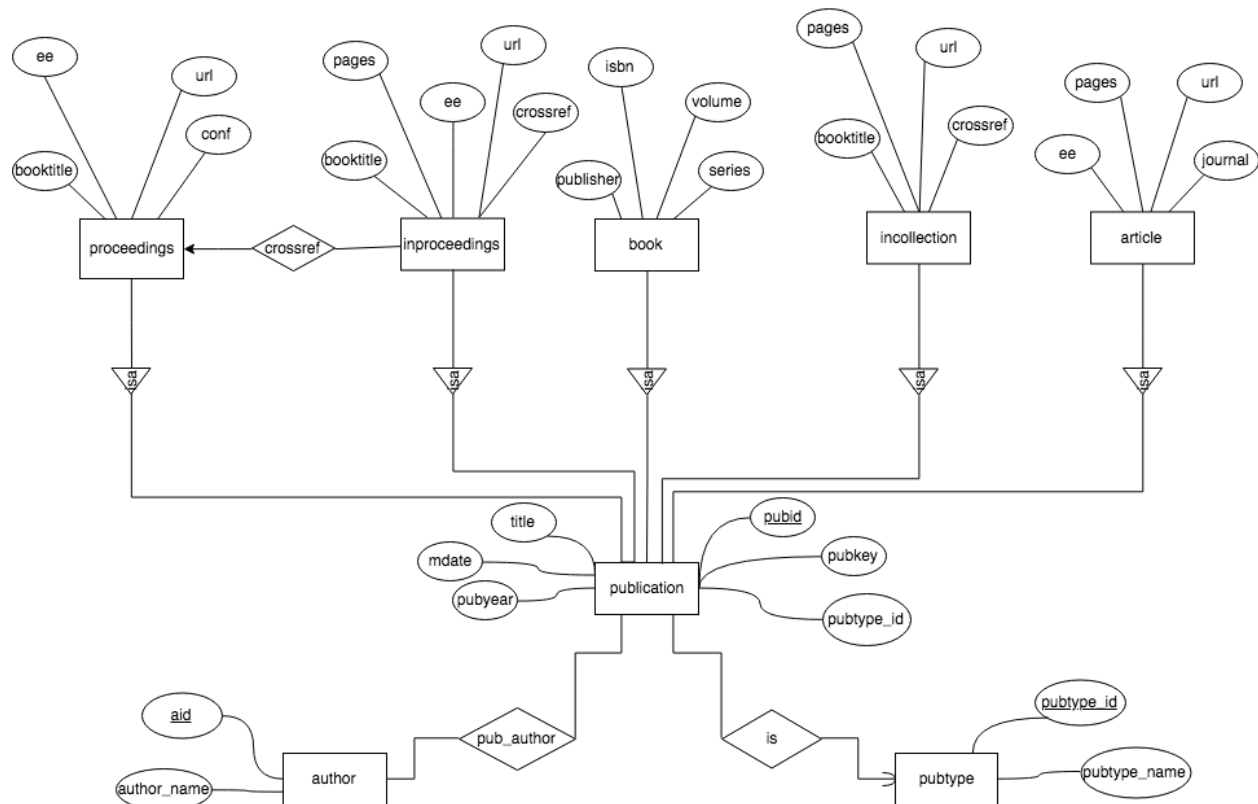


Figure 1. ER diagram

## 1.2 Commands for creating tables

*-- Table structure for table article*

```
CREATE TABLE article (  
  pubid    int(11) NOT NULL,  
  pages    varchar(255) DEFAULT NULL,  
  ee       varchar(255) DEFAULT NULL,  
  url      varchar(500) DEFAULT NULL,  
  journal  varchar(255) DEFAULT NULL,  
  PRIMARY KEY (pubid),  
  CONSTRAINT article_publication_pubid_fk FOREIGN KEY (pubid) REFERENCES  
publication (pubid)  
)  
  ENGINE = InnoDB  
  DEFAULT CHARSET = utf8mb4  
  COLLATE = utf8mb4_0900_ai_ci;
```

*-- Table structure for table author*

```
CREATE TABLE author (  
  aid      int(11) NOT NULL,  
  authorname varchar(255) DEFAULT NULL,  
  PRIMARY KEY (aid),  
  UNIQUE KEY author_authorname_uindex (authorname)  
)  
  ENGINE = InnoDB  
  DEFAULT CHARSET = utf8mb4  
  COLLATE = utf8mb4_0900_ai_ci;
```

*-- Table structure for table book*

```
CREATE TABLE book (  
  pubid    int(11) NOT NULL,  
  publisher varchar(255) DEFAULT NULL,  
  isbn     varchar(255) DEFAULT NULL,  
  series   varchar(255) DEFAULT NULL,  
  volume   varchar(255) DEFAULT NULL,  
  PRIMARY KEY (pubid),  
  CONSTRAINT book_publication_pubid_fk FOREIGN KEY (pubid) REFERENCES  
publication (pubid)  
)  
  ENGINE = InnoDB  
  DEFAULT CHARSET = utf8mb4  
  COLLATE = utf8mb4_0900_ai_ci;
```

*-- Table structure for table incollection*

```
CREATE TABLE incollection (  
  pubid    int(11) NOT NULL,  
  booktitle varchar(255) DEFAULT NULL,  
  pages     varchar(255) DEFAULT NULL,
```

```

    crossref  varchar(255) DEFAULT NULL,
    url       varchar(255) DEFAULT NULL,
    PRIMARY KEY (pubid),
    CONSTRAINT incollection_publication_pubid_fk FOREIGN KEY (pubid) REFERENCES
publication (pubid)
)
ENGINE = InnoDB
DEFAULT CHARSET = utf8mb4
COLLATE = utf8mb4_0900_ai_ci;

```

*-- Table structure for table inproceedings*

```

CREATE TABLE inproceedings (
    pubid      int(11) NOT NULL,
    booktitle  varchar(300) DEFAULT NULL,
    pages      varchar(255) DEFAULT NULL,
    crossref_pubid int(11)      DEFAULT NULL,
    ee         varchar(255) DEFAULT NULL,
    url        varchar(255) DEFAULT NULL,
    PRIMARY KEY (pubid),
    KEY inproceedings_crossref_pubid_fk (crossref_pubid),
    CONSTRAINT inproceedings_crossref_pubid_fk FOREIGN KEY (crossref_pubid)
REFERENCES proceedings (pubid)
)
ENGINE = InnoDB
DEFAULT CHARSET = utf8mb4
COLLATE = utf8mb4_0900_ai_ci;

```

*-- Table structure for table proceedings*

```

CREATE TABLE proceedings (
    pubid      int(11) NOT NULL,
    conf       varchar(255) DEFAULT NULL,
    ee         varchar(255) DEFAULT NULL,
    url        varchar(500) DEFAULT NULL,
    booktitle  varchar(255) DEFAULT NULL,
    PRIMARY KEY (pubid),
    CONSTRAINT proceedings_publication_pubid_fk FOREIGN KEY (pubid) REFERENCES
publication (pubid)
)
ENGINE = InnoDB
DEFAULT CHARSET = utf8mb4
COLLATE = utf8mb4_0900_ai_ci;

```

*-- Table structure for table pub\_author*

```

CREATE TABLE pub_author (
    pubid int(11) NOT NULL,
    aid   int(11) NOT NULL,
    PRIMARY KEY (pubid, aid),
    KEY pub_author_aid_fk (aid),
    CONSTRAINT pub_author_aid_fk FOREIGN KEY (aid) REFERENCES author (aid)
    CONSTRAINT pub_author_aid_fk FOREIGN KEY (aid) REFERENCES author (aid)
)

```

```

)
ENGINE = InnoDB
DEFAULT CHARSET = utf8mb4
COLLATE = utf8mb4_0900_ai_ci;

-- Table structure for table publication
CREATE TABLE publication (
  pubid      int(11) NOT NULL,
  pubkey     varchar(255) DEFAULT NULL,
  pubtype_id int(11)      DEFAULT NULL,
  title      varchar(1660) DEFAULT NULL,
  mdate      date          DEFAULT NULL,
  pubyear    year(4)       DEFAULT NULL,
  PRIMARY KEY (pubid),
  KEY publication_pubyear_index (pubyear),
  KEY publication_pubtype__fk (pubtype_id),
  CONSTRAINT publication_pubtype__fk FOREIGN KEY (pubtype_id) REFERENCES
pubtype (pubtype_id)
)
ENGINE = InnoDB
DEFAULT CHARSET = utf8mb4
COLLATE = utf8mb4_0900_ai_ci;

-- Table structure for table pubtype
CREATE TABLE pubtype (
  pubtype_id int(11) NOT NULL,
  pubname     varchar(255) DEFAULT NULL,
  PRIMARY KEY (pubtype_id)
)
ENGINE = InnoDB
DEFAULT CHARSET = utf8mb4
COLLATE = utf8mb4_0900_ai_ci;

```

## 1.3 Data Acquisition

\*Note: Please refer to Appendix for Java code for SAX parser.

We have created a SAX parser that is able to parse 'dblp.xml' into three files

1. publication.csv
2. author.csv
3. pub\_author.csv.

### 1.3.1 publication.csv

This file stores all 7 subclasses of publication. Each row is a publication (pub\_id = row number). The columns are all the possible fields for every publication (total 28 fields). If a publication does not have a certain field, that column is skipped.

pubid	pubtype	mdate	key	author	journal	book	booktitle	cdrom	chapter	crossref	cite	editor	ee
1	article	28/5/17	journals/ac	Sanjeev Saxena	Acta Inf.								<a href="https://doi.org/10.1007/BF03036466">https://doi.org/10.1007/BF03036466</a>
2	article	28/5/17	journals/ac	Hans Ulrich Simon	Acta Inf.								<a href="https://doi.org/10.1007/BF01257084">https://doi.org/10.1007/BF01257084</a>
3	article	28/5/17	journals/ac	Oded Shmueli	Acta Inf.								<a href="https://doi.org/10.1007/BF00289414">https://doi.org/10.1007/BF00289414</a>

Figure 2. Snapshot of 'publication.csv' file (28 columns are truncated due to lack of space)

### 1.3.2 author.csv

This file contains a list of authors. Each row has two columns: author\_id, and author\_name.

aid	name
1	Sanjeev Saxena
2	Hans Ulrich Simon
3	Nathan Goodman

Figure 3. Snapshot of 'author.csv' file

### 1.3.3 pub\_author.csv

This file has two columns: pub\_id and author\_name.

It stores a many-to-many relationship between publications and authors. For example, in Fig4, pubid 3 has 2 authors: Nathan Goodman and Oded Shmueli respectively.

pubid	author_name
1	Sanjeev Saxena
2	Hans Ulrich Simon
3	Nathan Goodman
3	Oded Shmueli

Figure 4. Snapshot of 'pub\_author.csv' file

### 1.3.4 Populate database tables

Process:

1. Import 3 '.csv' files into 3 temporary tables
  - a. publication.csv -> publication\_temp
  - b. author.csv -> author\_temp
  - c. pub\_author.csv -> pub\_author\_temp
2. Populate the 7 subclasses of publication by selecting necessary rows and columns from 'publication\_temp' table.

E.g:

```
INSERT INTO inproceedings (pubid, booktitle, pages, crossref, ee, url)
SELECT pub.pubid, pub.booktitle, pub.pages, pub.crossref, pub.ee,
pub.url
FROM publication
WHERE pubtype = 1; //inproceeding
```

3. Remove duplicate authors in 'author.csv'
4. Update pub\_author table by replacing author\_name with author\_id
5. Drop unnecessary columns in publication\_temp table and keep only the columns appeared in pubSchema.
6. We noticed that inproceeding references proceeding  
E.g :

```

<inproceedings key="conf/kdd/FayyadCRPCL17" mdate="2017-08-25">
  <author>Usama M. Fayyad</author>
  <author>Arno Candel</author>
  <author>Eduardo Ariño de la Rubia</author>
  <author>Szilárd Pafka</author>
  <author>Anthony Chong</author>
  <author>Jeong-Yoon Lee</author>
  <title>
    Benchmarks and Process Management in Data Science: Will We Ever
    Get Over the Mess?
  </title>
  <pages>31-32</pages>
  <year>2017</year>
  <booktitle>KDD</booktitle>
  <ee>https://doi.org/10.1145/3097983.3120998</ee>
  <crossref>conf/kdd/2017</crossref>
  <url>db/conf/kdd/kdd2017.html#FayyadCRPCL17</url>
</inproceedings>

```

```

<proceedings key="conf/kdd/2017" mdate="2017-08-15">
  <title>
    Proceedings of the 23rd ACM SIGKDD International Conference on
    Knowledge Discovery and Data Mining, Halifax, NS, Canada, August
    13 - 17, 2017
  </title>
  <booktitle>KDD</booktitle>
  <publisher>ACM</publisher>
  <year>2017</year>
  <isbn>978-1-4503-4887-4</isbn>
  <ee>http://doi.acm.org/10.1145/3097983</ee>
  <url>db/conf/kdd/kdd2017.html</url>
</proceedings>

```

we insert a column into inproceeding table so that it contains a crossref\_pubid which stores the pubid of its corresponding proceeding.

## 7. Drop temporary tables



## 2. Queries and Optimizing Queries

We have created the required queries and ran on MySQL. The output is captured in screenshot and the execution time is recorded. Next, we prepare 2 more dataset which contains 50% and 25% of original DBLP data. Using 3 databases with different sizes, we aim to find the effect of size of database on query time.

### 2.1 SQL queries & screen capture of results

*# (1) For each type of publication, count the total number of publications of that type between 2000-2017.*

*Query:*

```
WITH temp_count AS(  
    SELECT pubtype_id, COUNT(*) AS count  
    FROM publication  
    GROUP BY pubtype_id  
)  
  
SELECT pubtype.pubname, temp_count.count  
FROM temp_count JOIN pubtype  
WHERE temp_count.pubtype_id = pubtype.pubtype_id;
```

*Result:*

	pubname	count
1	article	1885783
2	inproceeding	2248769
3	proceeding	38248
4	book	15418
5	incollection	47420
6	phdthesis	67475
7	www	2153410

# (2) Find all the conferences that have ever published more than 200 papers in one year and are held in July.

Query:

```
WITH count_greater_than_200 AS (  
    SELECT crossref_pubid AS pubid, COUNT(*) AS count  
    FROM inproceedings  
    GROUP BY crossref_pubid  
    HAVING COUNT(*) > 200  
)  
  
conf_in_JULY AS (  
    SELECT pubid, pubkey, title  
    FROM publication  
    WHERE title LIKE "%JULY%" and pubtype_id = 2  
)  
  
required_conf_dup AS (  
    SELECT C1.pubid  
    FROM count_greater_than_200 AS C1  
    INNER JOIN conf_in_JULY AS C2 ON C1.pubid = C2.pubid  
)  
  
SELECT DISTINCT P.conf  
FROM proceedings AS P  
INNER JOIN required_conf_dup AS R ON P.pubid = R.pubid;
```

(result table is on the following page)

Result:

70 rows					
	conf				
1	icccn	23	icetet	47	ecis
2	icdar	24	issi	48	amcc
3	esiat	25	aaai	49	isbi
4	liss	26	IEEEcca	50	icml
5	acl	27	eusflat	51	icpads
6	iscc	28	dac	52	iceei
7	icorr	29	icdma	53	embc
8	cec	30	eurocon	54	icufn
9	csndsp	31	urai	55	iros
10	icnc	32	ni	56	icmcs
11	pacis	33	isspa	57	cvpr
12	tsp	34	dihu	58	cogsci
13	vcip	35	ecoopw	59	icls
14	ijcnn	36	fskd	60	ivs
15	fusion	37	isit	61	icalt
16	igarss	38	siggraph	62	gecco
17	cira	39	icita	63	trustcom
18	indin	40	iwcmc	64	fuzzIEEE
19	pdpta	41	IEEEcit	65	eucc
20	icdsp	42	chinasip	66	ecai
21	aied	43	sigir	67	atal
22	wce	44	aimech	68	iaiaiai
		45	snpd	69	ijcai
		46	compsac	70	icmlc

# (3a) Find the publications of author = "X" (Rolf Hennicker) at year 2015.

Query:

```
WITH temp_publication AS(
    SELECT P.*, A.authorname
    FROM (pub_author AS PA
    INNER JOIN publication AS P ON PA.pubid = P.pubid)
    INNER JOIN author AS A ON PA.aid = A.aid
)

SELECT *
FROM temp_publication
WHERE temp_publication.authorname = "Rolf Hennicker" and
temp_publication.pubyear = 2015;
```

Result:

	pubid	pubkey	pubtype_id	title	mdate	pubyear	authorname
1	1409	journals/acta/HennickerK15	0	Moving from interface theories to assembly theories.	2017-05-28	2015	Rolf Hennicker
2	1714950	journals/fac/MadeiraMBH15	0	Refinement in hybridised institutions.	2017-06-06	2015	Rolf Hennicker
3	2066384	conf/birthday/HennickerKW15	1	Model-Checking Helena Ensembles with Spin.	2017-05-23	2015	Rolf Hennicker
4	2066958	conf/birthday/NicolaH15	1	A Homage to Martin Wirsing.	2017-05-23	2015	Rolf Hennicker
5	2954553	conf/facs2/BelznerHW15	1	OnPlan: A Framework for Simulation-Based Online Planning.	2017-05-19	2015	Rolf Hennicker
6	6375761	series/lncs/MayerVKHPTPKB15	4	The Autonomic Cloud.	2017-05-16	2015	Rolf Hennicker

# (3b) Find the publications of author = "X" (Chun Tang) at year "Y" (2003) at conference "Z" (KDD).

Query:

```
WITH temp_conference AS(
    SELECT *
    FROM publication
    WHERE pubkey LIKE ("conf/kdd/%") and pubyear = 2003
)

SELECT *
FROM (temp_conference JOIN pub_author ON temp_conference.pubid =
pub_author.pubid)
JOIN author ON pub_author.aid = author.aid
WHERE author.authorname = "Chun Tang";
```

Result:

	pubid	pubkey	title	pubyear	aid	authorname
1	3354792	conf/kdd/TangZP03	Mining phenotypes and informative genes from gene expression data.	2003	39800	Chun Tang

# (3c). Find authors who published at least 2 papers at conference "Z" (KDD) at year "Y" (2003).

Query:

```
WITH temp_conference AS(
    SELECT *
    FROM publication
    WHERE pubkey LIKE ("conf/kdd/%") and pubyear = 2003
```

)

```
SELECT author.authurname
FROM (temp_conference JOIN pub_author ON temp_conference.pubid =
pub_author.pubid)
JOIN author ON pub_author.aid = author.aid
GROUP BY author.authurname
HAVING count(*) > 1;
```

*Result:*

	authurname
1	William DuMouchel
2	Aidong Zhang
3	Jian Pei
4	Dimitrios Gunopulos
5	Jiawei Han 0001
6	Jiong Yang
7	Inderjit S. Dhillon
8	Charu C. Aggarwal
9	David D. Jensen
10	Sheng Ma
11	Philip S. Yu
12	Bing Liu 0001
13	Ke Wang 0001
14	Mohammed Javeed Zaki
15	Padhraic Smyth
16	Xintao Wu
17	Yong Ye
18	William W. Cohen
19	Robert F. Murphy
20	Eamonn J. Keogh

# (4a). All authors who published at least 10 PVLDB papers and published at least 10 SIGMOD papers.

*Query:*

```
WITH temp_pub_author AS (
    SELECT pub_author.aid, pub_author.pubid, author.authurname
    FROM pub_author JOIN author ON pub_author.aid = author.aid
),
```

```
temp_pvladb AS (
```

```

SELECT temp_pub_author.aid, temp_pub_author.authorname, count(*)
FROM publication JOIN temp_pub_author ON publication.pubid =
temp_pub_author.pubid
WHERE pubkey LIKE ("conf/vldb/%")
GROUP BY temp_pub_author.aid, temp_pub_author.authorname
HAVING COUNT(*) > 9
),

```

```

temp_sigmod AS (
    SELECT temp_pub_author.aid, temp_pub_author.authorname, count(*)
    FROM publication JOIN temp_pub_author ON publication.pubid =
temp_pub_author.pubid
    WHERE pubkey LIKE ("conf/sigmod/%")
    GROUP BY temp_pub_author.aid, temp_pub_author.authorname
    HAVING COUNT(*) > 9
)

```

```

SELECT *
FROM temp_pvldb
WHERE temp_pvldb.authorname IN
(SELECT authorname
FROM temp_sigmod);

```

(result table is on the following page)

Result:

	aid	authorname	÷ 'count(*)' ÷		aid	authorname	÷ 'count(*)' ÷
1	37382	David J. DeWitt	29	23	31264	Kian-Lee Tan	15
2	31141	Christos Faloutsos	25	24	798156	Kenneth A. Ross	12
3	31145	Krithi Ramamritham	13	25	34803	Ioana Manolescu	10
4	31785	Raghu Ramakrishnan	30	26	88614	Daniela Florescu	12
5	32088	S. Sudarshan 0001	17	27	798058	Donald Kossmann	15
6	31040	Beng Chin Ooi	14	28	32398	Michael J. Franklin	22
7	2689	Carlo Zaniolo	11	29	79940	Serge Abiteboul	15
8	31519	Nick Roussopoulos	12	30	798650	Tova Milo	12
9	12999	Alfons Kemper	21	31	30934	Philip S. Yu	21
10	2350	Rakesh Agrawal 0001	27	32	31005	Jiawei Han 0001	21
11	8435	Umeshwar Dayal	20	33	12598	Miron Livny	13
12	1644	Philip A. Bernstein	21	34	31380	Michael J. Carey 0001	26
13	31255	Christian S. Jensen	13	35	15089	Hector Garcia-Molina	27
14	32501	Jennifer Widom	20	36	35257	Dan Suciu	12
15	31127	Divyakant Agrawal	16	37	33213	Stefano Ceri	13
16	798448	Stanley B. Zdonik	17	38	31624	Divesh Srivastava	23
17	36456	Yannis E. Ioannidis	20	39	35055	Yannis Papakonstanti...	10
18	31561	H. V. Jagadish	35	40	33673	Gerhard Weikum	25
19	32643	Nick Koudas	22	41	31006	Laks V. S. Lakshmanan	19
20	1151	Abraham Silberschatz	21	42	32085	Rajeev Rastogi	13
21	32945	Guy M. Lohman	14	43	31265	Yufei Tao	11
22	36264	Michael Stonebraker	25	44	430122	C. Mohan 0001	15
61 rows							
	aid	authorname	÷ 'count(*)' ÷		aid	authorname	÷ 'count(*)' ÷
40	33673	Gerhard Weikum	25				
41	31006	Laks V. S. Lakshmanan	19				
42	32085	Rajeev Rastogi	13				
43	31265	Yufei Tao	11				
44	430122	C. Mohan 0001	15				
45	34632	Martin L. Kersten	11				
46	39857	Wolfgang Lehner	16				
47	32949	Bruce G. Lindsay 0001	11				
48	39646	Volker Markl	13				
49	38888	Jeffrey F. Naughton	27				
50	32050	Minos N. Garofalakis	15				
51	36596	Joseph M. Hellerstein	11				
52	32729	Surajit Chaudhuri	26				
53	33337	Sihem Amer-Yahia	10				
54	37447	Dennis E. Shasha	13				
55	20386	Guido Moerkotte	14				
56	31643	Elke A. Rundensteiner	12				
57	2421	Henry F. Korth	10				
58	31235	Per-Åke Larson	13				
59	33925	Goetz Graefe	10				
60	31710	K. Selçuk Candan	10				
61	32998	Wenfei Fan	11				

# (4b). all authors who published at least 15 PVLDB papers but never published a KDD paper.

Query:

```
WITH temp_pub_author AS (  
    SELECT pub_author.aid, pub_author.pubid, author.authorname  
    FROM pub_author JOIN author ON pub_author.aid = author.aid  
    ),  
  
temp_pvldb AS (  
    SELECT temp_pub_author.aid, temp_pub_author.authorname, count(*)  
    FROM publication JOIN temp_pub_author ON publication.pubid =  
    temp_pub_author.pubid  
    WHERE pubkey LIKE ("conf/vldb/%")  
    GROUP BY temp_pub_author.aid, temp_pub_author.authorname  
    HAVING COUNT(*) > 14  
    ),  
  
temp_kdd AS (  
    SELECT DISTINCT temp_pub_author.aid, temp_pub_author.authorname  
    FROM publication JOIN temp_pub_author ON publication.pubid =  
    temp_pub_author.pubid  
    WHERE pubkey LIKE ("conf/kdd/%")  
    )  
  
SELECT *  
FROM temp_pvldb  
WHERE temp_pvldb.authorname NOT IN (  
    SELECT authorname  
    FROM temp_kdd);
```

Result:

	aid	authorname	count(*)
1	32088	S. Sudarshan 0001	17
2	12999	Alfons Kemper	21
3	1644	Philip A. Bernstein	21
4	798448	Stanley B. Zdonik	17
5	36456	Yannis E. Ioannidis	20
6	36264	Michael Stonebraker	25
7	32398	Michael J. Franklin	22
8	79940	Serge Abiteboul	15
9	31380	Michael J. Carey 0001	26
10	316099	Georges Gardarin	15
11	430122	C. Mohan 0001	15
12	38888	Jeffrey F. Naughton	27
13	12949	Patrick Valduriez	18
14	244975	Hasso Plattner	15

# (5). For each 10 consecutive years starting from 1970, compute the total number of conference publications in DBLP in that 10 years.

Query:

```
WITH year_temp AS(  
    SELECT DISTINCT P.pubyear, FLOOR((P.pubyear-1970)/10) AS decade_no  
    FROM pub AS P
```



```

WHERE P.pubyear >= 1970
ORDER BY P.pubyear ASC
),

pcount AS(
SELECT P.pubyear, count(P.pubid) AS psum
FROM pub AS P
WHERE P.pubkey LIKE ("conf/%") and P.pubyear >= 1970
GROUP BY P.pubyear
)

SELECT (decade_no * 10 + 1970) AS start, (decade_no * 10 + 1980) AS end,
SUM(pcount.psum) AS total
FROM pcount JOIN year_temp ON pcount.pubyear = year_temp.pubyear
GROUP BY decade_no
ORDER BY decade_no ASC;

```

*Result:*

	start ÷	end ÷	total ÷
1	1970	1980	12899
2	1980	1990	52971
3	1990	2000	222033
4	2000	2010	798909
5	2010	2020	1193632

# (6). Find the most collaborative authors who published in a conference or journal whose name contains “data”.

Query:

```
WITH pub_title_contain_data AS (  
    SELECT pubid, title  
    FROM publication  
    WHERE title LIKE "%DATA%"  
),  
  
pubid_author_count as (  
    SELECT P1.pubid, COUNT(*) AS num_coauthor  
    FROM pub_title_contain_data AS P1  
    INNER JOIN pub_author AS P2  
    ON P1.pubid = P2.pubid  
    GROUP BY P1.pubid  
),  
  
aid_coauthor_count AS (  
    SELECT P2.aid, P1.num_coauthor  
    FROM pubid_author_count AS P1 INNER JOIN pub_author AS P2  
    ON P1.pubid = P2.pubid  
),  
  
most_coauthor AS (  
    SELECT aid, SUM(num_coauthor - 1 ) AS total_coauthor  
    FROM aid_coauthor_count  
    GROUP BY aid  
    ORDER BY total_coauthor DESC  
    LIMIT 1  
)  
  
SELECT A2.authorname, A1.total_coauthor  
FROM most_coauthor AS A1 INNER JOIN author AS A2 ON A1.aid = A2.aid;
```

Result:

	authorname	÷	total_coauthor ÷
1	Jiawei Han 0001		781

# (7). Find the top 10 authors with the largest number of publications that are published in conferences and journals whose titles contain word "Data" in the last 5 years.

Query:

```
WITH conf_DATA AS (  
    SELECT P.pubid  
    FROM proceedings AS I  
    LEFT JOIN publication AS P ON I.pubid = P.pubid  
    WHERE P.title LIKE "%DATA%" and P.pubyear > 2013  
),  
  
inproceedings_in_conf AS (  
    SELECT I.pubid AS pubid  
    FROM conf_DATA AS C  
    INNER JOIN inproceedings AS I ON C.pubid = I.crossref_pubid  
),  
  
aid_count AS (  
    SELECT P.aid, COUNT(*) AS count  
    FROM inproceedings_in_conf AS I  
    INNER JOIN pub_author AS P ON I.pubid = P.pubid  
    GROUP BY P.aid  
    ORDER BY COUNT(*) DESC  
    LIMIT 10  
)  
  
SELECT A2.authurname, A1.count  
FROM aid_count AS A1 JOIN author AS A2 on A1.aid = A2.aid;
```

(result table is on the following page)

Result:

	authorname	count
1	Philip S. Yu	127
2	Wolfgang Lehner	94
3	Christos Faloutsos	75
4	Jiawei Han 0001	71
5	Charu C. Aggarwal	58
6	Hui Xiong	57
7	Enhong Chen	52
8	Xiaofang Zhou	51
9	Alfredo Cuzzocrea	51
10	Xuemin Lin	50

# (8). List the name of the conferences such that it has ever been held in June, and the corresponding proceedings contain more than 100 publications.

Query:

```
WITH count_greater_than_100 AS (  
    SELECT crossref_pubid AS pubid, COUNT(*) AS count  
    FROM inproceedings  
    GROUP BY crossref_pubid  
    HAVING COUNT(*) > 200  
),  
  
conf_in_june AS (  
    SELECT pubid, pubkey, title  
    FROM publication  
    WHERE title LIKE "%june%" and pubtype_id = 2  
),  
  
required_conf_dup AS (  
    SELECT C1.pubid  
    FROM count_greater_than_100 AS C1  
    INNER JOIN conf_in_june AS C2 ON C1.pubid = C2.pubid  
)  
  
SELECT DISTINCT P.conf  
FROM proceedings AS P INNER JOIN required_conf_dup AS R ON P.pubid = R.pubid;
```

(result table is on the following page)

Result:

	conf		conf		conf
1	iccs	21	pdpta	41	isbi
2	geoinformatics	22	mipro	42	icca
3	isie	23	issi	43	icml
4	acl	24	hpcc	44	pssc
5	ACISicis	25	eusflat	45	digra
6	iscc	26	dac	46	icmcs
7	vtc	27	icc	47	cvpr
8	icra	28	urai	48	icls
9	semeval	29	ni	49	ivs
10	wcnis	30	biorob	50	icassp
11	ascc	31	naacl	51	lrec
12	cec	32	dihu	52	gecco
13	pacis	33	ecoopw	53	trustcom
14	mim	34	isit	54	fuzzIEEE
15	ijcnn	35	iwcmc	55	eucc
16	iscas	36	IEEEcit	56	cso
17	ice-itmc	37	isnn	57	icse
18	fusion	38	ecis	58	icdcs
19	cars	39	ifsa	59	atal
20	spawc	40	amcc	60	ipps

# (9a). Find authors who have published at least 1 paper every year in the last 30 years, and whose family name start with 'H'.

Query:

```

WITH pub_last_30_years AS (
    SELECT *
    FROM publication
    WHERE pubyear BETWEEN 1988 AND 2018
),

author_fname_H AS (
    SELECT *
    FROM author
    WHERE SUBSTRING_INDEX(author.authorname , ' ', -1) LIKE 'H%'
),

author_H_pub AS (
    SELECT A.authorname, P.pubid
    FROM author_fname_H AS A JOIN pub_author AS P ON A.aid = P.aid
),

```

```
author_H_pub_30 AS (
    SELECT DISTINCT A.authorname, P.pubyear
    FROM author_H_pub AS A
    INNER JOIN pub_last_30_years AS P ON A.pubid = P.pubid
)
```

```
SELECT A.authorname
FROM author_H_pub_30 AS A
GROUP BY A.authorname
HAVING COUNT(*) >30;
```

Result:

32 rows	
	authorname
1	Alan R. Hevner
2	Ali R. Hurson
3	Amir Herzberg
4	David Harel
5	David Haussler
6	David Hutchison
7	Dorit S. Hochbaum
8	Eduard H. Hovy
9	Frank van Harmelen
10	Geoffrey E. Hinton
11	James A. Hendler
12	Jenq-Neng Hwang
13	John P. Hayes
14	Joseph Y. Halpern
15	Juraj Hromkovic
16	Manuel V. Hermenegildo
17	Mark Horowitz
18	Matthew Hennessy
19	Maurice Herlihy
20	Michael Hanus
21	Nicholas J. Higham
22	Pascal Van Hentenryck
23	Peter G. Harrison
24	Pierre Hansen
25	Richard I. Hartley
26	Rolf Hennicker
27	Scott E. Hudson
28	Seth Hutchinson
29	Theo Härder
30	Thomas S. Huang
31	Vincent Hayward
32	Wen-meï W. Hwu

# (9b) Find the names and number of publications for authors who have the earliest publication record in DBLP.

Query:

```
WITH earliest_pub as (  
    SELECT author.aid AS aid, author.authorname AS authorname  
    FROM publication, pub_author, author  
    WHERE publication.pubyear = (  
        SELECT MIN(pubyear)  
        FROM publication)  
    AND publication.pubid = pub_author.pubid  
    AND pub_author.aid = author.aid  
)  
  
SELECT authorname, COUNT(*) AS cnt  
FROM earliest_pub as t1, pub_author  
WHERE t1.aid = pub_author.aid  
GROUP BY t1.aid,authorname;
```

Result:

	authorname	cnt
1	Arnold F. Emch	8
2	Frederic Brenton Fitch	31
3	W. V. Quine	58
4	C. J. Ducasse	5
5	J. Barkley Rosser	38
6	C. I. Lewis	2
7	Alonzo Church	14
8	Emil L. Post	4

# (10). Find the top 3 authors with largest number of publications in DBLP.

Query:

```
WITH author_top3 AS (  
    SELECT aid, COUNT(*) AS cnt  
    FROM pub_author  
    GROUP BY aid  
    ORDER BY COUNT(*) DESC  
    LIMIT 3  
)  
  
SELECT A.authorname, author_top3.cnt  
FROM author_top3 INNER JOIN author AS A ON author_top3.aid = A.aid;
```

Result:

	authorname	cnt
1	H. Vincent Poor	1595
2	Mohamed-Slim Alouini	1212
3	Philip S. Yu	1170



## 2.2 Figures and analysis

Table 1 records the different execution time when we ran the same queries on full DBLP database, half of DBLP publication records and quarter of data respectively.

Table 1 shows that execution time is roughly proportional to data size. When data size is halved, the average execution time is 57.77%. When data size is a quarter of the original size, the execution time became 20.15% of the original execution time. This is because the larger the data size, the more number of blocks needed to store the data, leading to a more significant disk I/O time.

We also notice that the decrease in execution time is not linear (100% -> 57.77% -> 20.15%). When the data size becomes 25%, the average drop in execution time is more than 75%. This may be because the original data size cannot fit into memory and hence two pass algorithm is required whereas smaller data size might be able to fit into memory and only single pass algorithm is required. Hence, the number of disk I/O can be decreased significantly. (e.g 3a, 3b)

Furthermore, the improvement on execution time varies for different queries. For example, when data size is halved, the execution time for query 2 is 80.6% of the original time. On the other hand, the execution time for query 6 drops to 43.5%.

Query	Full / ms	Half / ms	Quarter / ms	Half vs Full	Quarter vs Full
1	2625	1287	569	49.03%	21.68%
2	1299	1047	530	80.60%	40.80%
3a	10608	8413	112	79.31%	1.06%
3b	2374	1114	53	46.93%	2.23%
3c	1860	883	780	47.47%	41.94%
4a	3396	1630	1067	48.00%	31.42%
4b	3390	1997	737	58.91%	21.74%
5	20437	10995	4584	53.80%	22.43%
6	8016	3486	1826	43.49%	22.78%
7	6983	3489	650	49.96%	9.31%
8	1388	524	210	37.75%	15.13%
9a	25892	15632	7957	60.37%	30.73%

9b	20390	20156	961	98.85%	4.71%
10	10155	5521	1637	54.37%	16.12%
			<b>AVERAGE</b>	<b>57.77%</b>	<b>20.15%</b>

Table 1. Execution time for different data sizes

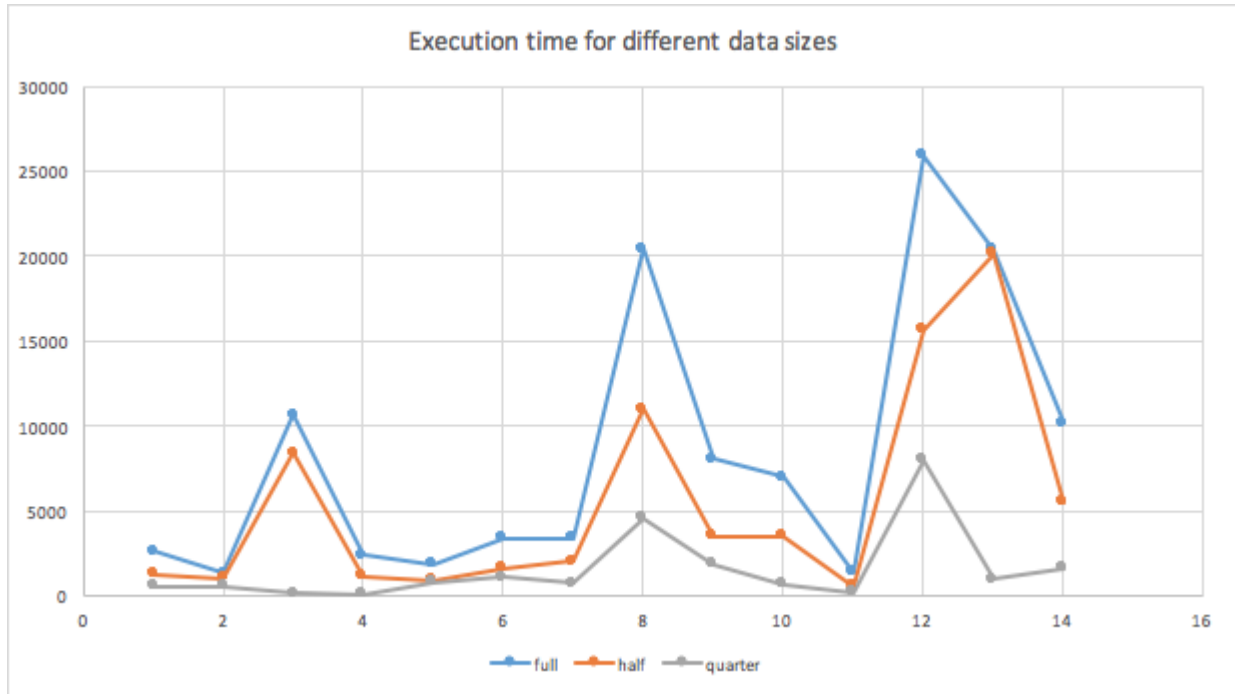


Figure 5. Execution time for different data sizes

## 3. Effect of Index

### 3.1 Create index command

```
CREATE UNIQUE INDEX author_authorname_uindex  
ON author (authorname);
```

```
CREATE UNIQUE INDEX inproceedings_crossref_pubid  
ON inproceedings (crossref_pubid);
```

```
CREATE UNIQUE INDEX pub_author_aid  
ON pub_author (aid);
```

```
CREATE INDEX publication_pubtype_id  
ON publication (pubtype_id);
```

```
CREATE INDEX publication_pubyear  
ON publication (pubyear);
```

```
CREATE INDEX publication_pubkey  
ON publication (pubkey);
```

### 3.2 Analysis on Index

Index is a file that accelerates the retrieval of records from a data file. Table 2 shows the execution time before and after adding the indices.


We noticed that the indices created in section 3.1 have different effects on different queries. For example, the decrease in execution time for query 6 is 18.73% whereas that for query 9b is 99.68%. This is because query 6 has a subquery “**WHERE title LIKE** “%DATA%”” and indices does not help in LIKE operator. On the other hand, query 9 has RANGE query on *pubyear*. Hence, an index on *pubyear* will improve execution time significantly for query 9.

Since *pubid* is unique for every publication and *aid* is unique for every author. Creating index for these attributes will significantly speed up processing queries which look for certain publication or author. As for the non-unique attributes, for example, *authorname* and *pubkey*, adding indices still plays an important role to upgrade the performance. Even though the indices cannot tell the placement of certain record, they provide information about the location of record and shorten the scanning time.

We also created an index on `inproceeding(crossref_pubid)`. This helps the performance when we want to group the inproceedings by the its proceeding pubid. This index is very useful for query 2, 7, 8 where we need to find all the inproceedings in a particular proceeding(conference).

Query	Before / ms	After / ms	Difference / ms	Change
1	2625	1423	1202	45.79%
2	1299	790	509	39.18%
3a	10608	129	10479	98.78%
3b	2374	265	2109	88.84%
3c	1860	663	1197	64.35%
4a	3396	168	3228	95.05%
4b	3390	119	3271	96.49%
5	20437	7912	12525	61.29%
6	8016	6515	1501	18.73%
7	6983	559	6424	91.99%
8	1388	690	698	50.29%
9a	25892	7879	18013	69.57%
9b	20390	66	20324	99.68%
10	10155	3575	6580	64.80%

*Table 2. Execution time before and after adding the index*



*Figure 5. Percentage change in execution time*



## 4. Effect of Cache


### 4.1 Analysis on effect of cache

In MySQL InnoDB storage engine, environment variable `innodb_buffer_pool_size` is used for configuring the size of buffer pool which is used for caching data and indexes in main memory.

In the experiments, we have compared the execution time for all 14 given queries with `innodb_buffer_pool_size` being 5GB, 512MB, 128MB, and 64MB, and the result has been shown in *Table 3* and *Figure 6*.

Query	5G / ms	512M / ms	128M / ms	64M / ms
1	926	940	1006	1020
2	581	713	757	719
3a	18	21	24	27
3b	25	24	25	17
3c	30	45	47	66
4a	79	140	124	115
4b	81	145	140	140
5	6204	7748	7576	6749
6	5952	6815	6975	6920
7	391	590	637	583
8	280	645	659	721
9a	21234	62264	98058	123378
9b	18	20	23	20
10	3161	3385	3674	3407

*Table 3. Execution time for different `innodb_buffer_pool_size` settings*



*Figure 6. Percentage change in execution time*

We could see that, the larger `innodb_buffer_pool_size` is, in general, the better the performance is. The reason being that is with larger pool size, more data or indexes could be cached in main memory, which leads to less disk I/O. Especially, taking query *9a* as an example, when the query is very complicated, the advantages of larger `innodb_buffer_pool_size` become more significant.

## 5. Appendix

### 5.1 Java code SAX Parser

```
import org.xml.sax.Attributes;
import org.xml.sax.SAXException;
import org.xml.sax.helpers.DefaultHandler;

import java.io.BufferedWriter;
import java.io.FileWriter;
import java.io.IOException;
import java.util.ArrayList;
import java.util.HashMap;
import java.util.List;

/**
 * This SAX parser is able to parse dblp.xml into three files: publictaion.csv, author.csv, and
 * pub_author.csv.
 * <ol>
 * <li>publication.csv</li>
 * This file stores all 7 subclasses of publication. Each row is a publication (pub_id = row
 * number)
 * The columns are all the possible fields for every publication.
 * If a publication does not have a certain field, that column is skipped.
 * <li>author.csv</li>
 * This file contains a list of authors. Each row has two columns: author_id, and author_name
 * <li>pub_author.csv</li>
 * This file has two columns: pub_id and author_name.
 * It stores a many-to-many relationship between publications and authors.
 * </ol>
 */
public class UserHandler extends DefaultHandler {
    private List<String> mPubElements;
    private List<String> mFieldElements;
    private HashMap<String, String> mValues;
    private int mPubId = 0; //unique pubid for each publication
    private int mAuthorId = 0; //unique author_id for each author
    private String mDelimiter = "\t";
    private BufferedWriter mWritePubAuthor;
    private BufferedWriter mWriterPub;
    private BufferedWriter mWriterAuthor;
    private StringBuilder mSb;

    public UserHandler() {
        mSb = new StringBuilder();

        //list of subclasses of publications
        mPubElements = new ArrayList<>();
        mPubElements.add("article");
        mPubElements.add("inproceedings");
        mPubElements.add("proceedings");
    }
}
```



```

mPubElements.add("book");
mPubElements.add("incollection");
mPubElements.add("phdthesis");
mPubElements.add("masterthesis");
mPubElements.add("www");

//list of all possible fields for a publication
mFieldElements = new ArrayList<>();
mFieldElements.add("pubid");
mFieldElements.add("pubtype");
mFieldElements.add("mdate");
mFieldElements.add("pubkey");
mFieldElements.add("booktitle");
mFieldElements.add("address");
mFieldElements.add("author");
mFieldElements.add("cdrom");
mFieldElements.add("chapter");
mFieldElements.add("city");
mFieldElements.add("crossref");
mFieldElements.add("cite");
mFieldElements.add("editor");
mFieldElements.add("ee");
mFieldElements.add("isbn");
mFieldElements.add("journal");
mFieldElements.add("month");
mFieldElements.add("note");
mFieldElements.add("number");
mFieldElements.add("pages");
mFieldElements.add("publisher");
mFieldElements.add("publnr");
mFieldElements.add("school");
mFieldElements.add("series");
mFieldElements.add("title");
mFieldElements.add("url");
mFieldElements.add("volume");
mFieldElements.add("year");

mValues = new HashMap<>();
}

```

```

/**
 * Handles differently a publication element from a field element.
 * If it is a publication element, all attributes are retrieved.
 * If it is a field element, there is no attribute.
 * @param uri
 * @param LocalName
 * @param qName
 * @param attributes
 * @throws SAXException
 */
@Override
public void startElement(String uri,
                        String localName,
                        String qName,
                        Attributes attributes) throws SAXException {
    qName = qName.toLowerCase();
    if (mPubElements.contains(qName)) {
        String mdate = attributes.getValue("mdate");
        String key = attributes.getValue("key");
        mValues.clear();
        mPubId++;
        mValues.put("pubid", String.valueOf(mPubId));
        mValues.put("pubtype", qName);
        mValues.put("mdate", mdate);
        mValues.put("pubkey", key);

    } else if (mFieldElements.contains(qName)) {
        if (qName.equalsIgnoreCase("author"))
            mAuthorId++;
    }
}

```

```

/**
 * Handles differently a publication element from a field element.
 * If it is the end of a publication element: insert one entry into "publication.csv"
 * If it is the end of a field element: insert into mValues which stores key-value pairs of
all fields.
 * A special case is 'author' element.
 * If it is the end element of 'author' element, we have to insert one entry into
'author.csv' and one entry into 'pub_author.csv'
 * @param uri
 * @param LocalName
 * @param qName
 * @throws SAXException
 */
@Override
public void endElement(String uri,
                      String localName, String qName) throws SAXException {

    qName = qName.toLowerCase();
    if (mPubElements.contains(qName)) {
        writeNewPub(mValues);
    } else if (mFieldElements.contains(qName)) {
        mValues.put(qName, new String(mSb));
        mSb.setLength(0);
        if (qName.equalsIgnoreCase("author")) {
            writeNewAuthor(mValues.get("author"));
            writeNewPubAuthor(mValues.get("author"));
        }
    }
}
}

```

```

/**
 * Read content between a start element and a end element.
 * Store the characters in a string builder.
 * Handles special characters such as "\"
 * @param ch
 * @param start
 * @param Length
 * @throws SAXException
 */
@Override
public void characters(char ch[], int start, int length) throws SAXException {
    String str = new String(ch, start, length);
    mSb.append(str.replaceAll("\\\\", "\\\\"));
}

/**
 * Insert one entry into 'publication.csv'
 * @param values
 */
private void writeNewPub(HashMap<String, String> values) {

    try {
        mWriterPub = new BufferedWriter(new FileWriter("publication.csv", true));
        for (String field : mFieldElements) {
            if (values.containsKey(field)) {
                if (mFieldElements.indexOf(field) != 0)
                    mWriterPub.write(mDelimiter);
                mWriterPub.write(values.get(field));
            } else {
                mWriterPub.write(mDelimiter);
            }
        }
        mWriterPub.write("\n");
        mWriterPub.close();
    } catch (IOException e) {
        e.printStackTrace();
    }
}

```

```
/**
 * Insert one entry into 'author.csv'
 * @param author
 */
private void writeNewAuthor(String author) {
    try {
        mWriterAuthor = new BufferedWriter(new FileWriter("author.csv", true));
        mWriterAuthor.write(mAuthorId + mDelimiter + author);
        mWriterAuthor.write("\n");
        mWriterAuthor.close();
    } catch (IOException e) {
        e.printStackTrace();
    }
}
```

```
/**
 * Insert one entry into 'pub_author.csv'
 * @param author
 */
private void writeNewPubAuthor(String author) {
    try {
        mWritePubAuthor = new BufferedWriter(new FileWriter("pub_author.csv", true));
        mWritePubAuthor.write(mPubId + mDelimiter + author);
        mWritePubAuthor.write("\n");
        mWritePubAuthor.close();
    } catch (IOException e) {
        e.printStackTrace();
    }
}
```